

LAB5

Name: Kushal Sourav B

Regno: 2347125

Displaying the Images

```
In [2]: import matplotlib.pyplot as plt
import os
from PIL import Image


def display_samples(dataset_dir, categories, num_samples=2):
    fig, axes = plt.subplots(len(categories), num_samples, figsize=(10, 10))
    for i, category in enumerate(categories):
        folder = os.path.join(dataset_dir, category)
        for j in range(num_samples):
            img_path = os.path.join(folder, os.listdir(folder)[j])
            img = Image.open(img_path)
            axes[i, j].imshow(img)
            axes[i, j].set_title(category)
            axes[i, j].axis('off')
    plt.tight_layout()
    plt.show()

categories = ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
display_samples('dataset/seg_train/seg_train/', categories)
```

buildings



buildings



forest



forest




glacier




glacier



mountain




mountain



sea




sea



street



street



Model Architecture:

```
In [3]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout, BatchNormalization

model = Sequential()

# 1st Convolutional Block
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(150, 150, 3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())

# 2nd Convolutional Block
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
```

```

model.add(BatchNormalization())

# 3rd Convolutional Block
model.add(Conv2D(128, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(BatchNormalization())

# Fully Connected Layers
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(6, activation='softmax'))

```

c:\Users\USER\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
 super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Model Training:

```

In [12]: from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Image Data Augmentation
train_datagen = ImageDataGenerator(rescale=1./255, validation_split=0.2) # 80/20 split

train_generator = train_datagen.flow_from_directory(
    'dataset/seg_train/seg_train/',
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='training')

validation_generator = train_datagen.flow_from_directory(
    'dataset/seg_train/seg_train/',
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='validation')

# Compile model
model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(train_generator, epochs=20, validation_data=validation_generator)

```

Found 11230 images belonging to 6 classes.

Found 2804 images belonging to 6 classes.

Epoch 1/20

c:\Users\USER\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
 self._warn_if_super_not_called()

```

351/351 ————— 315s 884ms/step - accuracy: 0.4831 - loss: 5.0357 - val_accuracy: 0.3531 - v
al_loss: 6.9323
Epoch 2/20
351/351 ————— 249s 709ms/step - accuracy: 0.5745 - loss: 1.4792 - val_accuracy: 0.6455 - v
al_loss: 1.2076
Epoch 3/20
351/351 ————— 234s 666ms/step - accuracy: 0.6516 - loss: 0.9739 - val_accuracy: 0.6740 - v
al_loss: 1.1260
Epoch 4/20
351/351 ————— 253s 721ms/step - accuracy: 0.7071 - loss: 0.8253 - val_accuracy: 0.7496 - v
al_loss: 0.7145
Epoch 5/20
351/351 ————— 245s 696ms/step - accuracy: 0.7432 - loss: 0.7367 - val_accuracy: 0.7568 - v
al_loss: 0.8130
Epoch 6/20
351/351 ————— 284s 810ms/step - accuracy: 0.7556 - loss: 0.6726 - val_accuracy: 0.7739 - v
al_loss: 0.7900
Epoch 7/20
351/351 ————— 263s 748ms/step - accuracy: 0.7719 - loss: 0.6267 - val_accuracy: 0.7846 - v
al_loss: 0.6722
Epoch 8/20
351/351 ————— 244s 694ms/step - accuracy: 0.7963 - loss: 0.5608 - val_accuracy: 0.7468 - v
al_loss: 0.7361
Epoch 9/20
351/351 ————— 222s 631ms/step - accuracy: 0.7963 - loss: 0.5657 - val_accuracy: 0.8067 - v
al_loss: 0.6560
Epoch 10/20
351/351 ————— 256s 730ms/step - accuracy: 0.8186 - loss: 0.5093 - val_accuracy: 0.7885 - v
al_loss: 0.6749
Epoch 11/20
351/351 ————— 269s 765ms/step - accuracy: 0.8267 - loss: 0.4883 - val_accuracy: 0.7846 - v
al_loss: 0.7140
Epoch 12/20
351/351 ————— 245s 696ms/step - accuracy: 0.8368 - loss: 0.4558 - val_accuracy: 0.7186 - v
al_loss: 0.9805
Epoch 13/20
351/351 ————— 232s 661ms/step - accuracy: 0.8580 - loss: 0.4046 - val_accuracy: 0.6837 - v
al_loss: 1.1641
Epoch 14/20
351/351 ————— 229s 651ms/step - accuracy: 0.8643 - loss: 0.3843 - val_accuracy: 0.8046 - v
al_loss: 0.6261
Epoch 15/20
351/351 ————— 229s 651ms/step - accuracy: 0.8858 - loss: 0.3146 - val_accuracy: 0.8167 - v
al_loss: 0.8053
Epoch 16/20
351/351 ————— 252s 718ms/step - accuracy: 0.8823 - loss: 0.3204 - val_accuracy: 0.7447 - v
al_loss: 0.8746
Epoch 17/20
351/351 ————— 256s 730ms/step - accuracy: 0.8819 - loss: 0.3324 - val_accuracy: 0.6762 - v
al_loss: 2.3130
Epoch 18/20
351/351 ————— 254s 723ms/step - accuracy: 0.8986 - loss: 0.2931 - val_accuracy: 0.8074 - v
al_loss: 0.7799
Epoch 19/20
351/351 ————— 251s 715ms/step - accuracy: 0.9047 - loss: 0.2619 - val_accuracy: 0.7657 - v
al_loss: 1.2976
Epoch 20/20
351/351 ————— 246s 699ms/step - accuracy: 0.9135 - loss: 0.2453 - val_accuracy: 0.7842 - v
al_loss: 0.9805

```

Evaluation:

```

In [13]: import matplotlib.pyplot as plt

def plot_history(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))

    plt.figure(figsize=(10,5))

```

```


plt.subplot(1,2,1)
plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

plot_history(history)

```



```


In [15]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np

Y_pred = model.predict(validation_generator)
y_pred = np.argmax(Y_pred, axis=1)

cm = confusion_matrix(validation_generator.classes, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d')
plt.title('Confusion Matrix')
plt.show()

```

88/88 ————— 13s 147ms/step



Optimization

```
In [4]: from tensorflow.keras.optimizers import Adam
from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest',
    validation_split=0.2)

train_generator = train_datagen.flow_from_directory(
    'dataset/seg_train/seg_train/',
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='training')

validation_generator = train_datagen.flow_from_directory(
    'dataset/seg_train/seg_train/',
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='validation')

# Compile model
model.compile(optimizer=Adam(), loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(train_generator, epochs=20, validation_data=validation_generator)
```

Found 11230 images belonging to 6 classes.

Found 2804 images belonging to 6 classes.

Epoch 1/20

```
c:\Users\USER\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\trainers\data_adapters
\py_dataset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)`
in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pas
s these arguments to `fit()`, as they will be ignored.
```

```
self._warn_if_super_not_called()
351/351 ————— 315s 885ms/step - accuracy: 0.3856 - loss: 5.5625 - val_accuracy: 0.3566 - v
al_loss: 3.0143
Epoch 2/20
351/351 ————— 246s 699ms/step - accuracy: 0.4612 - loss: 1.4952 - val_accuracy: 0.5357 - v
al_loss: 1.2523
Epoch 3/20
351/351 ————— 254s 720ms/step - accuracy: 0.5257 - loss: 1.2428 - val_accuracy: 0.4722 - v
al_loss: 1.6088
Epoch 4/20
351/351 ————— 270s 768ms/step - accuracy: 0.5670 - loss: 1.1505 - val_accuracy: 0.6205 - v
al_loss: 0.9796
Epoch 5/20
351/351 ————— 254s 722ms/step - accuracy: 0.6081 - loss: 1.0669 - val_accuracy: 0.5068 - v
al_loss: 1.2464
Epoch 6/20
351/351 ————— 254s 721ms/step - accuracy: 0.6156 - loss: 1.0243 - val_accuracy: 0.4900 - v
al_loss: 1.5121
Epoch 7/20
351/351 ————— 269s 763ms/step - accuracy: 0.6279 - loss: 1.0115 - val_accuracy: 0.5849 - v
al_loss: 1.0594
Epoch 8/20
351/351 ————— 263s 746ms/step - accuracy: 0.6491 - loss: 0.9474 - val_accuracy: 0.6255 - v
al_loss: 1.1502
Epoch 9/20
351/351 ————— 262s 744ms/step - accuracy: 0.6644 - loss: 0.9183 - val_accuracy: 0.6391 - v
al_loss: 0.9254
Epoch 10/20
351/351 ————— 249s 707ms/step - accuracy: 0.6788 - loss: 0.8754 - val_accuracy: 0.6926 - v
al_loss: 0.9011
Epoch 11/20
351/351 ————— 248s 705ms/step - accuracy: 0.6898 - loss: 0.8524 - val_accuracy: 0.7133 - v
al_loss: 0.9001
Epoch 12/20
351/351 ————— 244s 692ms/step - accuracy: 0.7086 - loss: 0.8287 - val_accuracy: 0.7172 - v
al_loss: 0.7833
Epoch 13/20
351/351 ————— 244s 692ms/step - accuracy: 0.7179 - loss: 0.8030 - val_accuracy: 0.7539 - v
al_loss: 0.7564
Epoch 14/20
351/351 ————— 244s 693ms/step - accuracy: 0.7176 - loss: 0.7811 - val_accuracy: 0.6287 - v
al_loss: 1.2154
Epoch 15/20
351/351 ————— 248s 705ms/step - accuracy: 0.7328 - loss: 0.7229 - val_accuracy: 0.6958 - v
al_loss: 0.8219
Epoch 16/20
351/351 ————— 245s 695ms/step - accuracy: 0.7360 - loss: 0.7356 - val_accuracy: 0.6844 - v
al_loss: 0.8748
Epoch 17/20
351/351 ————— 244s 694ms/step - accuracy: 0.7521 - loss: 0.7032 - val_accuracy: 0.7354 - v
al_loss: 0.7952
Epoch 18/20
351/351 ————— 245s 695ms/step - accuracy: 0.7387 - loss: 0.7260 - val_accuracy: 0.7133 - v
al_loss: 0.8937
Epoch 19/20
351/351 ————— 244s 693ms/step - accuracy: 0.7589 - loss: 0.6831 - val_accuracy: 0.7439 - v
al_loss: 0.8052
Epoch 20/20
351/351 ————— 244s 694ms/step - accuracy: 0.7592 - loss: 0.6826 - val_accuracy: 0.6997 - v
al_loss: 0.8637
```

```
In [5]: import matplotlib.pyplot as plt

def plot_history(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))
```


```
plt.figure(figsize=(10,5))

plt.subplot(1,2,1)
plt.plot(epochs, acc, 'r', label='Training accuracy')
plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
plt.title('Training and validation accuracy')
plt.legend()

plt.subplot(1,2,2)
plt.plot(epochs, loss, 'r', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

plot_history(history)
```



```
In [6]: from sklearn.metrics import confusion_matrix
import seaborn as sns
import numpy as np

Y_pred = model.predict(validation_generator)
y_pred = np.argmax(Y_pred, axis=1)

cm = confusion_matrix(validation_generator.classes, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d')
plt.title('Confusion Matrix')
plt.show()
```

88/88 ————— 19s 216ms/step

